Experience Report: The Next 1100 Haskell Programmers

Jasmin Christian Blanchette   Lars Hupel   Tobias Nipkow   Lars Noschinski   Dmitriy Traytel

Fakultät für Informatik, Technische Universität München, Germany

Abstract

We report on our experience teaching a Haskell-based functional programming course to over 1100 students for two winter terms. The syllabus was organized around selected material from various sources. Throughout the terms, we emphasized correctness through QuickCheck tests and proofs by induction. The submission architecture was coupled with automatic testing, giving students the possibility to correct mistakes before the deadline. To motivate the students, we complemented the weekly assignments with an informal competition and gave away trophies in a award ceremony.

1.   Introduction

This paper reports on a mandatory Haskell-based functional programming course at the Technische Universität München. In the first iteration (winter semester of 2012–2013), there were 619 students enrolled. In the following winter semester (2013–2014), there were 553 students enrolled. The course ran for 15 weeks with one 90-minute lecture and one 90-minute tutorial each week. The weekly homework was graded, but the final grade was primarily determined by the examination. To make the homework more attractive, we coupled it with an informal programming competition.

The departmental course description does not prescribe a specific functional language but focuses on functional programming in general. In the previous two years, the course had been based on Standard ML. We have a strong ML background ourselves but chose Haskell because of its simple syntax, large user community, real-world appeal, variety of textbooks, and availability of QuickCheck [3]. The one feature we could well have done without is lazy evaluation; in fact, we wondered whether it would get in the way.

The course was mandatory for computer science (*Informatik*) and information systems (*Wirtschaftsinformatik*) students. All had learned Java in their first semester. The computer science students had also taken courses on algorithms and data structures, discrete mathematics, and linear algebra. The information systems students had only had a basic calculus course and were taking discrete mathematics in parallel.

The *dramatis personae* in addition to the students were lecturer Tobias Nipkow, who designed the course, produced the slides,[1] and gave the lectures; *Masters of TAs* Lars Noschinski and Lars Hupel, who directed a dozen teaching assistants (TAs) and took care of the overall organization; furthermore the (*Co*)*Masters of Competition* Jasmin Blanchette (*MC*) and Dmitriy Traytel (*CoMC*), who selected competition problems and ranked the solutions.

2.   Syllabus

The second iteration covered the following topics in order. (The first iteration was similar, with a few exceptions discussed below.) Each topic was the subject of one 90-minute lecture unless otherwise specified.

1. Introduction to functional programming [0.5 lecture]
2. Basic Haskell: `Bool`, QuickCheck, `Integer` and `Int`, guarded equations, recursion on numbers, `Char`, `String`, tuples
3. Lists: list comprehension, polymorphism, a glimpse of the Prelude, basic type classes (Num, `Eq`, `Ord`), pattern matching, recursion on lists (including accumulating parameters and non-primitive recursion); scoping rules by example [1.5 lectures]
4. Proof by structural induction on lists
5. Higher-order functions: `map`, `filter`, `foldr`, λ-abstractions, extensionality, currying, more Prelude [2 lectures]
6. Type classes [0.5 lecture]
7. Algebraic datatypes: `data` by example, the general case, Boolean formula case study, structural induction [1.5 lectures]
8. I/O, including files and web
9. Modules: module syntax, data abstraction, correctness proofs
10. Case study: Huffman coding
11. Lazy evaluation and infinite lists
12. Complexity and optimization
13. Case study: parser combinators

Most topics were presented together with examples or smaller case studies, of which we have only mentioned Boolean formulas. Moreover, two topics kept on recurring: tests (using QuickCheck) and proofs (by induction).

From day one, examples and case studies in class were accompanied by properties suitable for QuickCheck. Rather than concentrate all inductive proofs in the lecture about induction, we distributed them over the entire course and appealed to them whenever it was appropriate. A typical example: In a case study, a function is first defined via `map myg . map myf` and then optimized to `map (myg . myf)`, justified by a proof of `map` $(g \ . \ f)$ = `map` g . `map` f.

Much of the above material is uncontroversial and part of any Haskell introduction, but some choices deserve some discussion.

---

[1] `www21.in.tum.de/teaching/info2/WS1314/slides.pdf`

Induction. Against our expectations, induction was well understood, as the examination confirmed (Section 5). What may have helped is that we gave the students a rigid template for inductions. We went as far as requiring them to prove equations $l_1 = r_1$ not by one long chain of equalities $l_1 = \cdots = l_m = r_n = \cdots = r_1$ but by two reductions $l_1 = \cdots = l_m$ and $r_1 = \cdots = r_n$. This avoids the strange effect of having to shift to reverse gear halfway through the proof of $l_1 = r_1$. It must be stressed that we considered only structural induction, that we generally did not expect the students to think up auxiliary lemmas themselves, and that apart from extensionality and induction all reasoning was purely equational.

In Haskell, there is the additional complication that proofs by structural induction establish the property only for finite objects. Some authors restrict the scope of their lemmas to finite lists of defined elements [15], while others prove `reverse (reverse xs)` `= xs` without mentioning that it does not hold for partial or infinite lists [7]. Although some authors discuss finite partial objects and infinite objects [6, 15], we avoided them in our course—undefinedness alone is a can of worms that we did not want to open. Hence, we restricted ourselves to a total subset of Haskell in which "fast and loose reasoning" [4] is sound.

Input/output and monads. In the first iteration, I/O was covered toward the end of the course because it is connected with the advanced topic of monads. For much of the course, many students may have had the impression that Haskell is only a glorified pocket calculator. Therefore we moved I/O to an earlier point in the course. We also dropped monads, since the majority had not grasped them. This was not entirely satisfactory, because a minority of students had been excited by monads and expressed their disappointment.

Abstraction functions. In the lecture on modules and data abstraction, we also showed how to prove correctness of data representations (e.g., the representation of sets by lists). This requires an abstraction function from the representation back to the abstract type that must commute with all operations on the type. As the corresponding homework showed, we failed to convey this. In retrospect, it is outside the core functional programming syllabus, which is why it is absent from all the textbooks. The topic still appeared briefly in the lecture in the second iteration, but without exercises.

Laziness. Haskell's lazy evaluation strategy and infinite objects played only a very minor role and were introduced only toward the end. Initially, we were worried that laziness might confuse students when they accidentally stumble across it before it has been introduced, but this was not reported as a problem by any of the TAs. However, we could not give the students a good operational model of the language without laziness: All they knew initially was that equations were applied in some unspecified order. Even after we had explained laziness, it remained unclear to many students how exactly to determine what needs to be evaluated.

Complexity and optimization. Complexity considerations are seriously complicated by laziness. We found that the book by Bird [1] offered the best explanation. For time complexity, he notes that assuming eager evaluation is easier and still gives an upper bound. Therefore, we simply replaced lazy by eager evaluation for this lecture. The principles then applied to most programming languages, and one can cover key optimizations such as tail recursion.

Parser combinators. In the first iteration of the course, the lecture on parser combinators had been given about two thirds into the course, but many students had failed to grasp it. As a result we moved it to the end of the course for the second iteration. This means that it could not be covered by an exercise sheet and we have no hard feedback on how well the parser combinators were understood. It should be noted that this is the first time during their studies that the students are exposed to the technicalities of parsing.

Figure 1. Status page with exercise points (where blue bars denote the student's points and black markers denote the median of all students)

Figure 2. Homework submissions relative to the number of enrolled students

3. Exercises

Each week we released an exercise sheet with group and homework assignments. The main objective of the homework was to have the students actually program in Haskell. The submission infrastructure periodically ran automatic tests, giving the students fast feedback and an opportunity to correct mistakes before the deadline.

3.1 Assignments

A typical assignment sheet contained between three and five group exercises and about as many homework exercises. The group exercises were solved in 90-minute tutorial groups. There were 25 or 26 such groups, each with up to 24 students. Each exercise focused on a specific concept from the week's lecture. Many were programming exercises, but some required the students to write QuickCheck tests, carry out proofs, or infer an expression's type.

The homework assignments, to be solved individually, covered the same topics in more depth, sometimes in combination. They were optional, but in the first iteration of the course, the students who collected at least 40% of the possible points were awarded a bonus of 0.3 to the final grade, on a scale from 1.0 ($\approx$ A$^+$) to 5.0 ($\approx$ F). In the end, 281 students claimed the bonus. Furthermore, the (Co)MCs nominated one of the assignments to count as part of the competition (Section 4).

Much to our regret, there were several unpleasing and time-consuming incidents with plagiarism (Section 3.4). Thus, we decided to drop the bonus system in the second iteration of the course. As a countermeasure against the anticipated decrease in the number of homework submissions, we tried to motivate the students by providing a graphical overview of their homework grades in comparison with the median value of all submissions on the web (Figure 1) and ensured quick grading of the homeworks by the TAs. Still, the decrease was severe (Figure 2): In the first iteration, 75% of the enrolled students submitted the first homework; the number dropped and finally stayed below 40% after sheet 10. In the second iteration, it started with 50% and stayed below 20% after sheet 8.

Most of the exercises were well understood by those who did them, perhaps as they conformed closely to the lectures. A few important exceptions are noted below.

A group problem consisted of registering the polymorphic function type `a -> b` as an instance of the `Num` type class, so that $(f + g)$ `x == f x + g x` and similarly for the other operations. Many students did not understand what their task was, or why one would register

functions as numbers; and even those who understood the question had to realize that `b` must be an instance of `Num` and fight the problem's higher-order nature. We had more success two weeks later when we redid the exercise for a `Fraction` datatype and gently explained why it makes sense to view fractions as numbers.

Less surprisingly, many students had issues with λ-abstractions. They tended to use λs correctly with `map` and `filter` (although many preferred list comprehensions when given the choice), but other exercises revealed the limits of their understanding. One exercise required implementing a function `fixpoint` $eq\ f\ x$ that repeatedly applies f to x until $f^{n+1}\ x$ `'eq'` $f^n\ x$ and then using this function to solve concrete problems. Another exercise featured a deterministic finite automaton represented as a tuple, where the δ component is represented by a Haskell function.

One difficulty we continually faced when designing exercises is that the Internet provides too many answers. This was an issue especially in the first few weeks, when little syntax has been introduced. We did our best to come up with fresh ideas and, failing that, obfuscated some old ideas.

3.2 Submission and Testing Infrastructure

The university provides a central system for managing student submissions, but we built our own infrastructure so that we could couple it with automatic testing. Our submission system combines standard Unix tools and custom scripts. The students were given a secure shell (`ssh`) account on the submission server. They had to upload their submissions following a simple naming convention. The system generated test reports every 15 minutes using Quick-Check. Many students appear to have improved their submissions iteratively based on the system's feedback. The final reports were made available to the TAs but had no direct effect on grading.

To increase the likelihood that the submissions compile with the testing system, we provided a correctly named template file for each assignment, including the necessary module declarations and stub definitions f `= undefined` for the functions to implement. Nonetheless, many students had problems with the naming scheme (there are surprisingly many ways to spell "exercise"), causing their submissions to be ignored. These problems went away after we started providing a per-student graphical web page listing the status of all their assignments and announced a stricter grading policy.

A few exercises required writing QuickCheck properties for a function described textually. These properties had to take the function under test as argument, so that we could check them against secret reference implementations. Since higher-order arguments had not yet been introduced, we disguised the argument type using a type synonym and put the boilerplate in the template file.

The test reports included the compilation status, the result of each test, and enough information about the failed tests to identify the errors. The tests themselves were not revealed, since they often contained hints for a correct implementation. In cases where the input of the test case did not coincide with the input of the tested function, we had to explain this in the description or provide more details using QuickCheck's `printTestCase` function. Some care was needed because the function under test can throw exceptions, which are not caught by QuickCheck because of the lazy evaluation of `printTestCase`'s argument. We used the `Control.Spoon` package to suppress these exceptions.

To make the output more informative, we introduced an operator `==?` that compares the expected and actual results and reports mismatches using `printTestCase`.

We did not find any fully satisfactory way to handle very slow and nonterminating functions. QuickCheck's `within` combinator fails if a single test iteration takes too long, but these failures are confusing for correct code. Instead, we limited the test process's runtime, potentially leaving students with a truncated report.

3.3 Test Design

As regular users of the Isabelle proof assistant [10], we had a lot of experience with Isabelle's version of QuickCheck [2]. The tool is run automatically on each conjectured lemma as it is entered by the user to exhibit flaws, either in the lemma itself or in the underlying specification (generally a functional–logic program). Typically, the lemmas arise naturally as part of the formalization effort and are not designed to reveal bugs in the specification.

We designed our Haskell tests to expose the most likely bugs and capture the main properties of the function under test. We usually also included a test against a reference implementation. We soon found out that many bugs escaped the test suite because the Haskell QuickCheck's default setup is much less exhaustive than its Isabelle namesake's. For example, the Haskell random generator tends to produce much larger integers than the Isabelle one; as a result, random lists of integers rarely contain duplicates, which are essential to test some classes of functions. Worse, for polymorphic functions we did not realize immediately that type variables are instantiated with the unit type `()` by default (a peculiar choice to say the least). In contrast, Isabelle's version of QuickCheck supports random testing, exhaustive testing (cf. SmallCheck [12]), and narrowing (cf. Lazy SmallCheck [12], Agsy [9]), the default number of iterations is 250, and type variables are instantiated by small types. The differences between the two QuickCheck versions became painfully obvious with the competition exercises, as we will see in Section 4.

Following these initial difficulties, the Masters of TAs were appointed Masters of Tests and put in charge of setting up the testing framework properly. They immediately increased QuickCheck's number of iterations, decreased the maximum size parameter, and regained control by defining custom generators and instantiating type variables with small types. They also started using Small-Check to reliably catch bugs exposed by small counterexamples.

3.4 Plagiarism Detection

We considered it important to detect and deter plagiarism in the first year, both because individual bonuses should be earned individually and because learning functional programming requires doing some programming on one's own. Our policy was clear: Plagiarism led to forfeiture of the bonus for all involved parties.

To identify plagiarists, we used Moss [13] extended with a custom shell script to visualize the results with Graphviz [5]. The resulting graph connects pairs of submissions with simi-

Figure 3. Plagiarism graph excerpt featuring the Pastebin clique

lar features, with thicker edges for stronger similarities. Figure 3 shows an anonymized excerpt of the output for week 3.

A noteworthy instance of unintended sharing is the complete subgraph of thick edges in the middle of Figure 3. One of the involved students has used Pastebin (`http://pastebin.com/`) for his own purposes, without realizing that it would be indexed by Google and picked up by other students.

Moss's results are imprecise, with many false positives, so they must be analyzed carefully. Functional programming often allows short, canonical solutions. Unusual naming conventions, spacing, or bugs are useful clues. One could have thought that the recent German plagiarism scandals, which eventually cost two federal ministers their Dr. title and minister position, would have cured the country for some time. Sadly, we had to disqualify 29 students.

4. Competition

Our main inspiration for the programming competition has been CADE's Automated Theorem Prover System Competition [14], organized by Geoff Sutcliffe since 1996. We have been entering Isabelle since 2009 and have noticed the competition's impact on the theorem proving community. We were also moved by our late colleague Piotr Rudnicki's arguments in favor of contests [11]:

> I am dismayed by the watering down of the curriculum at CS departments which does not push the students to their intellectual limits. This wastes a lot of talented people who, under these conditions, have no chance to discover how talented and capable they really are. The programming contests attract a substantial fraction of the most talented students that we have; I enjoy working with them and they seem to enjoy doing it too.
>
> The Heavenly Father, with his unique sense of humor, has distributed the mental talents in a rather unpredictable way. It is our role to discover these talents and make them shine. If we do not do it, then we—the educators—will end up in Hell. And I would rather not get there just for this one reason.

For our own contest, each week we selected one of the programming assignments as a competition problem. We also fixed a criterion for ranking the correct entries. By enclosing their solutions within special tags, students became competitors. Each week, rank $i \in \{1, \ldots, 20\}$ brought in $21 - i$ points. The five students cumulating the most points were promised "tasteful" trophies.

Once the entries had been tested and ranked, we published the names of the top 20 students on the competitions' web pages[2] and updated the cumulative top 20. To avoid legal issues regarding privacy, we inserted a notice in the assignment sheets, making it clear that the competition is an opt-in. The list of winners was followed by a discussion of the most remarkable solutions, written by the MC, the CoMC, or the Masters of TAs, in ironic self-important third-person style.

An unexpected side effect of the competition is that it provided a channel to introduce more advanced concepts, such as higher-order functions, before they were seen in class. The criteria were designed to raise the students' awareness of engineering trade-offs, including performance and scalability, even though these topics were beyond the scope of the course.

As is to be expected, participation went down as the session progressed. We tended to lose those students who were not in the cumulative top 20, which is the reason why we extended it to a top 30 in the second iteration. The optional exercises attracted only the hard core. We have the testimony of a student who, after gathering enough points to secure the grade bonus, skipped the mandatory exercises to focus on the competition. Thus, our experience corroborates Rudnicki's: Contests motivate talented students, who otherwise might not get the stimuli they need to perform.

The (Co)MCs revealed the last competition week's results in an award ceremony during the last lecture, handing out the trophies and presenting characteristic code snippets from each winner.

Because of the various ranking criteria, and also because students knew that their solutions could turn up on the web page, the competition triggered much more diversity than usual assignments. Ranking the solutions was fascinating. Each week, we had plenty of material for the web page. The page was read by many students, including some who were not even taking the course. A summary of selected competition problems and solutions follows, originating from the first iteration of the course unless indicated otherwise.

---

[2] www21.in.tum.de/teaching/info2/WS1213/wettbewerb.html, www21.in.tum.de/teaching/info2/WS1314/wettbewerb.html

Week 1: Sum of Two Maxima's Squares (434 Entrants)

Task: Write a function that adds the squares of the two largest of its arguments x, y, z. *Criterion:* Token count (lower is better).

The task and criterion were modeled after a similar Scheme exercise by Jacques Haguel. By having to keep the token count low, students are encouraged to focus on the general case.

The winner's solution had 13 tokens (excluding the left-hand side and counting 'max' as one):

```
max x y ^ 2 + min x y 'max' z ^ 2
```

Here the concision was attained at the expense of simplicity, to the point that we felt the need to verify the solution with Isabelle. Lists appeared in several of the top 20 solutions:

```
sum $ tail $ (^2) 'map' sort [x, y, z]
sum [x * x | x <- tail (sort [x, y, z])]
```

A few competitors exploited the explicit application operator $ to save on parentheses ($f \$ g \ x$ == $f \ (g \ x)$). Using only syntaxes and functions seen in class, a 25-token solution was possible:

```
x * x + y * y + z * z - a * a
  where a = min x (min y z)
```

The median solution had a 3-way case distinction. There were plenty of 6-way distinctions, and one entry even featured a correct 10-way distinction using < and ==, complete with 64 needless parentheses, totaling 248 tokens. This provided the ideal context for the MC to quote Donald Knuth [8, p. 56] on the competition's web site: "The ability to handle lots of cases is Computer Science's strength and weakness. We are good at dealing with such complexity, but we sometimes don't try for unity when there is unity."

To count tokens, we initially used a fast-and-frugal Perl script. However, many students asked us to make the program available, so we replaced it by a bullet-proof Haskell solution based on Niklas Broberg's lexical analyzer (Language.Haskell.Exts.Lexer).

Week 5: Quasi-subsequences (206 Entrants)

Task: Write a function that tests whether a list $[x_1, \ldots, x_m]$ is a quasi-subsequence of a list ys, meaning that it is a (noncontiguous) subsequence of ys or that there exists an index k such that $[x_1, \ldots, x_{k-1}, x_{k+1}, \ldots, x_m]$ is a subsequence of ys. *Criterion:* Speed.

Thomas Genet shared this example with us. The problem statement mentioned that the MC's solution took 0.4 s for quasiSubseq [1 .. N] ([N] ++ [2 .. N − 1] ++ [1]) with $N = 50\,000$.

To rank the solutions, we ran some additional QuickCheck tests, filtering out 43 incorrect solutions from the 143 that compiled and passed all the official tests. Then we tried various examples, including the one above with different Ns, and eliminated solutions that reached the generous timeout. Some examples had a huge m and a short ys. This produced 20 winners, whom we listed on the web site. The algorithms varied greatly and were difficult to understand. One of the TAs, Manuel Eberl, contributed an automaton-based solution. The MC's solution had a dynamic programming flavor.

The story does not end here. Having noticed more bugs while ranking, we speculated that some of the top 20 entries might be incorrect. Prompted to action by Meta-Master Nipkow, we rechecked all 20 solutions using Isabelle's implementation of QuickCheck and found flaws in 6 of them. We did not penalize their authors but took a second look at Haskell's testing capabilities (cf. Section 3.2).

Week 6: Email Address Anonymizer (163 Entrants)

Task: Write a function that replaces all email addresses in a text by an anonymized version (e.g., p__.q___@f_____.c__). *Criterion:* Closeness to the official definition of email address.

The task idea came from Koen Claessen. The statement suggested a simple definition of email addresses, which is what most

students implemented, but pointed to RFCs 5321 and 5322 for a more precise definition. Our goal was to see how students react to unclear requirements. Of course, the RFCs partly contradicted each other, and it was not clear whether the domains had to be validated against the rules specified in earlier RFCs. It was also left to the student's imagination how to locate email addresses in the text.

This task was, by far, the most despised by the students. It was also the most difficult to rank to be fair to those who invested many hours in it but failed some simple test we had designed to rank the solutions. We revised the ranks upward to comfort the more vocal participants, turning the top 20 into a top 45.

Weeks 8–9: Boolean Solver (Optional, 14 Entrants)

Task: Write a Boolean satisfiability (SAT) solver. *Criterion:* Speed.

To avoid repeating the week 6 debacle, we suggested five optimizations to a DPLL algorithm that would be evaluated in turn:

1. Eliminate pure positive variables.
2. Select short clauses before long ones.
3. Select frequent literals before infrequent ones.
4. Use a dedicated algorithm for Horn formulas.
5. Use a dedicated algorithm for 2CNF.

Obvious variants of these optimizations would be invoked to break ties. The Meta-Master promised a Ph.D. degree for polynomial solutions (to no avail).

For the evaluation, we needed to devise problems that can be solved fast if and only if the heuristic is implemented. Showing the absence of an optimization turned out to be much more difficult than we had anticipated, because the various optimizations interact in complicated ways, and the exact mixture varied from solution to solution. To make matters worse, often the optimizations were implemented in a naive way that slowed down the solver (e.g., reprocessing the entire problem each time a literal is selected to detect whether an optimization has become applicable).

Unlike for previous weeks, this problem gave no homework points. In exchange, it was worth double (40 points), and the students had two weeks to complete it. Also, we did not provide any QuickCheck tests, leaving it to the students to think up their own ("just like in real life"). There were 14 submissions to rank, as well two noncompetitive entries by TA Eberl and one by the (Co)MCs. We found bugs in 8 of the submissions, but gave these incorrect solutions some consolation points to help populate the week's "top 20." The two best solutions implemented optimizations 1 to 4 and pure negative literal elimination, but neither 2CNF nor dual-Horn.

Week 13: Programmatic Art (Optional, 10 Entrants)

Task: Write a program that generates a pixel or vector graphic in some standard format. *Criteria:* Aesthetics and technique.

There were no constraints concerning the subject of the picture or its generation. While the imprecision had annoyed the students in the email anonymizer assignment, here it was perceived as a most welcome artistic freedom. The creations' visual nature was a perfect match for the award ceremony, where the weekly and final results were presented.

The students were asked to turn in both the program and a generated picture. The Meta-Master and the Masters of TAs rated the aesthetics on a scale from 1 to 10. The re-

maining 10 points were issued by the (Co)MCs for "technique," mostly as a safeguard against cheaters.

Two students "misunderstood" the exercise: one handed in a generated ASCII-art text file, another used `Network.Curl.Download` to retrieve Vincent van Gogh's "Starry Night" from Google. The latter secured the top score for aesthetics but was punished with a 2 for technique. The winner had been visibly inspired by Piet Mondrian's famous "Compositions" (Figure 4). His randomized solution could generate arbitrarily many fake Mondrians.

We ran this challenge again in the second iteration. Interestingly enough, a fair share of competitors produced animations, whereas in the first iteration, only static images were submitted. Both times, fractals (like the Mandelbrot set) were popular among the students.

### Weeks 11–12 (2nd iter.): Othello (Optional, 18 Entrants)

Task: Write a player based on artificial intelligence for Othello. *Criterion:* Ranking in a tournament between all contestants.

According to the students, this was one of the most time-consuming tasks, or—from a different viewpoint—one of the most enjoyable tasks to invest time in. The task was inspired by Tjark Weber. The contestants were allowed to turn in up to three different programs. Most provided only one implementation, but alternated some parameters (e.g., different values of depth for the game tree search). This was also the first and only exercise where students had to adhere to a specified module interface, where we specified some abstract types. There were 36 working programs in total. We ran all possible pairs of them with a timeout of 20 minutes per player per game. The evaluation of all 1260 games in the tournament took a few days on a cluster of 8 dual-core workstations. We were surprised to see that there was little to no correlation between the total running time of each player and their ranking—there were both extremely fast and good solutions (the third place obtained 90% of the points of the second place but took only 2% of the time), as well as slow and bad solutions.

5. Examination

The final grades were based on a two-hour examination. Students were allowed to bring one hand-written "cheat sheet." They needed 40% to pass. The results were then translated to a 1.0 to 5.0 scale. The 0.3 homework bonus was awarded only to those who had passed the examination. No bonus was applied in the second iteration. The correlation between homework points and examination points is shown in Figure 5.

The problems were similar to the group exercises but avoided more advanced or mundane topics (e.g., modules and data abstraction). The examination was designed so that the best students could finish in one hour. Perhaps because we had no previous experience in teaching Haskell, the marking revealed many surprises. Our impressions are summarized below for seven problems:

Figure 5. Correlation between points in the exercises and in the examination (second iteration)

Figure 4. The winning "Mondrian" entry

1. *Infer the types of given expressions.* Many students who understood types and type inference in practice had problems applying their knowledge in a more abstract context. They often forgot to instantiate type variables or to remove the argument type when applying a function. For example `filter not :: [Bool] -> [Bool]` was often typed as `[a] -> [a]` or even `(a -> Bool) -> [a] -> [a]`. Tellingly, one of the best students lost 2.5 of 5 points here, while answering all the other questions correctly.

2. *Implement the same simple function using recursion, using a list comprehension, and using higher-order functions (e.g.,* `map`*,* `filter`*).* The definitions based on a list comprehension were usually correct. The corresponding `map`–`filter` version proved more challenging. The recursive definitions were usually correct but sometimes lacked cases.

3. *Implement a function that lifts a variable renaming function to a logical formula datatype.* The datatype featured both a directly recursive constructor (for logical negation) and recursive constructors through lists (for *n*-ary conjunctions and disjunctions). The recursion through the list, using `map (rename f)`, confused many (although it had been covered in several exercises). Some solutions managed to change the shape of the formula, rewriting *n*-ary expressions into nested binary expressions. The pattern matching syntax was also not universally understood, and the constructors were often missing in the right-hand sides.

4. *Prove* `map` *f* `(concat` *xss*`) = concat (map (map` *f*`)` *xss*`).` The proof by induction posed little problems. Presumably the students had the induction template on their cheat sheet. Quite a few followed the template too slavishly, claiming to be doing an induction on *xs* instead of *xss*. Another common mistake was to take *xss* = `[[]]` as the base case.

5. *Choose two test functions from a given set that together constitute a complete test suite for a given function operating on lists.* There were seven tests to choose from: tests for the `[]`, `[`*x*`]`, and *x* : *xs* cases, a distributivity law, a length law, and two properties about the result list's content. Obvious solutions were `[]` with *x* : *xs* or `[`*x*`]` with distributivity, but there were many other combinations, most of which we discovered while marking. For example, the length law implies the `[]` case, and the `[`*x*`]` and *x* : *xs* cases together imply the `[]` case. Of all people, we should have been alert to the dangers of axiomatic specifications. It would have been easy to use Isabelle to prove or refute each of the 21 possible answers before sending the examination to press.

6. *Evaluate given expressions step by step.* The order of evaluation was not understood by all. We were downright shocked by some of the answers provided for `(\x -> (\y -> (1 + 2) + x)) 4 5`. We were not prepared to see monstrosities such as `(\4 -> (\5 -> (1 + 2) + 4))` as the end result.

7. *Write an I/O program that reads the user's input line by line and prints the total number of vowels seen so far.* The monadic solutions were surprisingly good, perhaps due to the students' familiarity with imperative programming. The main difficulty was to keep track of the cumulative vowel count. Many solutions simply printed the count of each line instead. Another common mistake was to use the monadic syntax `<-` instead of `let` to bind nonmonadic values.

Some statistics: 552 (493 in the second iteration) students registered for the exams. 432 (379) students took the examination. 334 (262) passed it. 39 (11) secured the top grade (1.0), with at least 47 out of 50 (38.5 out of 40) points. Five (one) had a perfect score. The higher results in the first iteration can be explained by the bonus, which undoubtedly led to more motivation for doing homework and better preparation for the examination. In total, this amounted to a difference in passing rate of 8 percentage points.

6. Conclusion

Teaching functional programming using Haskell has been an enjoyable experience overall. As is usually the case for functional programming, the feedback from students was mixed. If we have failed to convince some of them of the value of functional programming, we have also received many testimonies of students who have "seen the light," and some of the serious competitors told us the course had been the most fun so far.

For future years, we plan to either leave out some of the more advanced material or enhance its presentation. Type inference is one topic we downplayed so far; it should be possible to present it more rigorously without needing inference trees. On the infrastructure side, we developed tool support for simple proofs by induction, in the form of a lightweight proof assistant and will integrate it into our homework submission system in the upcoming third iteration.

References

[1] R. Bird. *Introduction to Functional Programming using Haskell*. Prentice Hall, 1998. Second edition.

[2] L. Bulwahn. The new Quickcheck for Isabelle—Random, exhaustive and symbolic testing under one roof. In C. Hawblitzel and D. Miller, editors, *CPP 2012*, volume 7679 of *LNCS*, pages 92–108. Springer, 2012.

[3] K. Claessen and J. Hughes. QuickCheck: A lightweight tool for random testing of Haskell programs. In *ICFP '00*, pages 268–279. ACM, 2000.

[4] N. A. Danielsson, J. Hughes, P. Jansson, and J. Gibbons. Fast and loose reasoning is morally correct. In J. G. Morrisett and S. L. P. Jones, editors, *POPL 2006*, pages 206–217. ACM, 2006.

[5] J. Ellson, E. R. Gansner, E. Koutsofios, S. C. North, and G. Woodhull. Graphviz—Open source graph drawing tools. In *Graph Drawing*, pages 483–484, 2001.

[6] P. Hudak. *The Haskell School of Expression*. Cambridge University Press, 2000.

[7] G. Hutton. *Programming in Haskell*. Cambridge University Press, 2007.

[8] D. E. Knuth, T. L. Larrabee, and P. M. Roberts. *Mathematical Writing*. Mathematical Association of America, 1989.

[9] F. Lindblad. Property directed generation of first-order test data. In M. Morazán, editor, *TFP 2007*, pages 105–123. Intellect, 2008.

[10] T. Nipkow, L. C. Paulson, and M. Wenzel. *Isabelle/HOL: A Proof Assistant for Higher-Order Logic*, volume 2283 of *LNCS*. Springer, 2002.

[11] P. Rudnicki. Why I do it? `http://webdocs.cs.ualberta.ca/~piotr/ProgContest/why.txt`. Accessed 25 Feb. 2013.

[12] C. Runciman, M. Naylor, and F. Lindblad. SmallCheck and Lazy SmallCheck: Automatic exhaustive testing for small values. In A. Gill, editor, *Haskell 2008*, pages 37–48. ACM, 2008.

[13] S. Schleimer, D. S. Wilkerson, and A. Aiken. Winnowing: Local algorithms for document fingerprinting. In A. Y. Halevy, Z. G. Ives, and A. Doan, editors, *SIGMOD Conference*, pages 76–85. ACM, 2003.

[14] G. Sutcliffe and C. B. Suttner. The state of CASC. *AI Commun.*, 19 (1):35–48, 2006.

[15] S. Thompson. *Haskell, the craft of functional programming*. Addison Wesley, 2011. Third edition.